

Review-to-Execution Continuity for Corrected Agent State Changes

Liang Hanghao Xuan Wentao Chen Qile Peng Peng

College of Computer Science and Electronic Engineering, Hunan University

Changsha 410082, China; corresponding author: hnu16pp@hnu.edu.cn

Abstract

State-changing agents can continue acting after a reviewer corrects and authorizes a machine proposal. An executed candidate may then have the same value and retained context as the reviewed candidate while being a different instance. We define correction-aware review-to-execution continuity by separating machine proposal, authorized correction, reviewed candidate and executed transition. An executable contract connects these roles to predecessor freshness, atomic successors, evidence integrity and field-level provenance. Bounded relational checks and independently audited persistence support the contract; two exact-binding realizations agree on 165 controlled cases. An audit recovers the correction relation in 84 archived hosted-agent chains. A locally frozen prospective experiment records 512 planned policy trajectories over 128 tasks and two deployments, retaining runtime failures and unknown integrity. Instance-bound execution prevents all 17 observed substitution violations in standardized queued handoffs. Actual rejection feedback induces reviewed-candidate reuse in 16 continuations and reauthorization in one. Agent-managed refreshed previews produce no observed substitutions. Overall completion changes by 0.78 percentage points, with a task-cluster 95% interval of $[-1.95, 3.52]$. The main deployment’s standardized handoffs require one additional agent tool call and median 14.54 seconds of suffix time. The other deployment reaches only 11 review checkpoints because of provider failures. These results expose continuity’s execution consequences while separating task utility, integrity and operational friction. **Keywords:** trustworthy agents; review-to-execution continuity; human authorization; transactional execution; provenance; agent recovery

1 Introduction

Agents that invoke tools can change operational state: they can update a record, publish a revised value, or commit a result that later actions treat as authoritative. The execution problem therefore extends beyond generating a plausible answer. A persistent change needs an accountable connection to the proposal and authorization that justified it. Human-in-the-loop frameworks already support pausing tool calls for approval, editing their arguments, and resuming execution [1]. These capabilities make the interval between review and later execution a concrete part of an agent workflow.

Correction makes that interval especially consequential. A machine proposes a value x_c , and a reviewer authorizes x_a . A valid correction permits $x_c \neq_{\text{json}} x_a$: if the machine reads 100 and the reviewer corrects it to 101, the authoritative state should contain 101 while preserving the original proposal. The proposal explains what the machine produced; the authorization explains why a different value may be committed. Guided repair and interactive cleaning already distinguish system suggestions from human judgments [2, 3]. Executing the corrected action adds a further question: which proposal instance was reviewed, and which instance actually supplied the committed transition?

Consider two persisted candidates c_1 and c_2 . Both propose 100 for the same field, refer to the same evidence and predecessor, and retain the same producer and review context. They differ only in instance identity. The reviewer inspects c_1 and authorizes 101. Before committing, the agent regenerates an equivalent preview or receives a queued handoff carrying c_2 . The final value may still be 101 and all contextual checks may agree. Yet review of c_1 does not authorize c_2 under an instance-accountability requirement. Equality of content and continuity of the reviewed instance answer different questions. A workflow must make that distinction explicit before deciding whether the continuation may mutate authoritative state.

Several established mechanisms can implement parts of the required connection. Transactions install atomic state changes [4]; concurrency checks bind them to a current predecessor [5]; provenance records derivation and responsibility [6, 7]; approval and request binding control which actions proceed [8, 9]. The question we examine is their joint correction case: how the machine proposal, authorized correction, reviewed candidate and executed transition remain distinguishable across persistence and subsequent agent actions. It has both a systems dimension—specifying and enforcing the relation—and an execution dimension—what the agent does when that relation produces rejection feedback.

We make correction-aware review-to-execution continuity explicit, enforceable, and observable in state-changing agent workflows. Our central question is: *How should authorization preserve review-to-execution continuity when corrected agent proposals become persistent changes, and how does enforcing that continuity reshape subsequent execution and recovery?* We compare retained-context and instance-bound authorization after the same prospective reviewed history,

tracing feedback to the next meaningful agent action and eventual utility and integrity outcomes. This links the persistence contract to the environment in which an adaptive agent continues its task.

The paper has four contributions:

1. **Correction-aware continuity.** We distinguish proposal value, authorized correction, reviewed instance and executed instance, and characterize the information needed to distinguish the declared continuity failures.
2. **Executable authorization contract.** We connect the reviewed target to freshness, complete successors, evidence and field sources, with bounded formal checking, transactional realizations and independent persisted-state auditing.
3. **Cross-layer mechanism evidence.** Controlled policy separation, request-binding tests and archived hosted-agent correction chains establish complementary evidence for the specified relation.
4. **Prospective online intervention.** A frozen paired experiment inserts the policies into live agent continuations to measure recovery paths, separate utility/integrity outcomes and operational friction.

Three findings organize the evidence. Correction creates a distinct continuity problem. The relation can be specified and enforced across persistence mechanisms. In standardized online handoffs, its enforcement redirects subsequent agent actions and adds measurable execution work. Seventeen context-admitted substitutions violate the instance-continuity oracle; the paired bound continuations instead complete with intact continuity. In the agent-managed preview tasks, neither policy produces an observed substitution. This contrast connects an execution-layer distinction to actual feedback and recovery, rather than presuming that agents ordinarily choose an invalid continuation.

2 Related Work

2.1 Agent approval and tool execution

Approval mechanisms already control state-changing operations. Clark–Wilson combines authenticated execution, well-formed transactions and reconstructible audit records [8]. Collaborative databases can retain cell versions pending approval, including approval tied to a cell identifier and timestamp [10]. Current agent middleware interrupts selected tool calls and permits approval, edits or rejection before resumption [1]. Server-controlled authorization data and a final execution gate are established security guidance [9]. Our correction case concerns the relation among four roles after such a decision: the original proposal, the corrected authorized value, the reviewed candidate and the candidate used for execution. A completed approval step provides a starting point for that relation; the continuation determines whether it is preserved.

2.2 Transactions, provenance and governed state

Transactions and optimistic concurrency control support atomic successors and validation against current state [4, 5]. Database provenance identifies contributing data and derivations [11, 12]. Curated-database provenance records edits and copies across versions, including the connection of unchanged data to preceding versions [6]; reenactment reconstructs transaction provenance [13]. PROV-DM separates entities, activities and agents [7]. These mechanisms supply the persistence vocabulary for the execution contract. The authorization target still has to specify whether equivalent contents or a particular reviewed instance may supply a corrected transition. We implement the same contract with relational objects and an event journal so that this question is separated from storage representation.

2.3 Trustworthy execution and recovery

AgentDojo evaluates task utility and security outcomes separately in interactive tool environments [14]. That separation motivates reporting successful task completion independently of the lineage of authoritative effects. Alloy and Kodkod make relational constraints executable within finite scopes [15, 16], while TestEra connects specifications to executions through abstraction and concretization [17]. A close neighboring preprint, Continuity Kernel, connects exact proposal identity, a predecessor, approver evidence and lineage to atomic accepted state.¹ Exact proposal binding is an established neighboring principle. Our distinction retains the machine proposal when a different value is authorized, ties that correction to the reviewed/executed instance, and observes the consequences of enforcement for live continuation. The studies discussed above establish approval, persistence, provenance and interactive evaluation capabilities. Our joint question isolates review-to-execution instance continuity under correction and observes what its enforcement makes an agent do next.

3 Review-to-Execution Continuity

3.1 Four roles in a corrected state change

Let Σ_v be an authoritative record state with a nonempty field domain F_r , total values $V_v : F_r \rightarrow X$ and total durable sources $S_v : F_r \rightarrow T$. A persisted candidate is

$$c = \langle id, r, f, e, v_{\text{exp}}, x_c, p \rangle, \quad (1)$$

¹He J, Yu D (2026), *Beyond Memory: A Transactional Continuity Kernel for Long-Lived AI Agents*, arXiv:2608.11632. Preprint discussed in text; its status and source comparison are retained in the supplement. <https://arxiv.org/abs/2608.11632>.




Figure 1: Correction-aware review-to-execution continuity. Review preserves the machine proposal but authorizes a corrected value. A content-equivalent new candidate separates the two policies: context admits it, whereas bound requires the reviewed instance. Rejection can lead the continuing agent to reuse that instance or obtain fresh authorization. The diagram is a motivating construction; observed continuation counts appear in Section 7.

where id names the immutable instance, r, f identify the target, e identifies persisted evidence and its locator, v_{exp} is the expected predecessor, and p identifies the producer. Retained context $K(c)$ includes the declared non-identity review information. Instance identity is not reconstructed from the proposal value or $K(c)$.

An authorization $a = \langle c_r, x_a, q \rangle$ records the reviewed instance c_r , explicit authorized value x_a and authorized principal q , together with its validity information. A commit request submits c_s and a value x . The roles obey

$$\begin{aligned} \text{proposal}(c_r) &= x_c, \\ \text{authorized}(a) &= x_a, \\ \text{committed}(\Sigma_{v+1}, f) &=_{\text{json}} x_a. \end{aligned} \quad (2)$$

The committed transition retains the proposal as origin evidence and the authorization as the authority for the chosen value. Correction permits $x_c \neq_{\text{json}} x_a$; it does not rewrite the immutable proposal to make the two values appear equal.

Contract value comparisons use canonical JSON equality, written $=_{\text{json}}$: sorted-key, whitespace-minimal serialization must agree. Consequently, an integer, a floating-point serialization and a Boolean remain distinguishable when their serialized values differ. This equality is shared by change detection, admission value checks, tracing and the independent oracle.

3.2 Content agreement and reviewed-instance continuity

For an already issued authorization, define

$$\begin{aligned} K_{eq} &\iff K(c_s) = K(c_r), \\ I_{eq} &\iff id(c_s) = id(c_r). \end{aligned} \quad (3)$$

Let G collect the remaining disclosed checks: authorization validity, predecessor freshness, committed-value

equality, evidence integrity, successor shape, source relations, transaction constraints and grant consumption. It excludes an implicit instance comparison. The two execution policies are

$$P_{\text{context}} : \text{Admit} = G \wedge K_{eq}, \quad (4)$$

$$P_{\text{bound}} : \text{Admit} = G \wedge K_{eq} \wedge I_{eq}. \quad (5)$$

Both retain the reviewed identifier. Only the decision to require its equality at execution differs. Context-equivalent c_2 can therefore pass Eq. 4 under c_1 's authorization while failing Eq. 5. This is a policy separator under a declared instance-accountability requirement. Context-based authorization remains a coherent choice for an application whose responsibility contract treats those instances as interchangeable.

3.3 Failure-distinguishing information

The declared failure model separates five information classes: candidate identity/context (D_C), proposal/authorized-value attribution (D_V), predecessor freshness (D_F), complete atomic successor (D_B), and total field sources (D_S). For a history h , $\pi_I(h)$ retains the specified logical observations in I , and $O(h)$ assigns the full-history authority outcome, including successor and trace status.

Proposition: class-wise irredundancy. For every $d \in \mathcal{D} = \{D_C, D_V, D_F, D_B, D_S\}$, the declared model contains paired histories satisfying

$$\pi_{\mathcal{D} \setminus \{d\}}(h_s) = \pi_{\mathcal{D} \setminus \{d\}}(h_u), \quad O(h_s) \neq O(h_u). \quad (6)$$

Thus a decision based on the reduced observations cannot distinguish both required outcomes.

Construction. Hold the other observations fixed while changing, respectively, candidate/context binding, the authority assigned to a corrected value, the

entry-time predecessor comparison, the single-versus-fragmented commit sequence, or one required source anchor. Each pair changes the corresponding derived authority outcome. A fixed-registry pair further selects either the reviewed candidate or a context-equivalent alternative while retaining the same two candidates and non-identity information. It isolates instance continuity within D_C . Observation functions and valid histories are specified independently; the checker derives projections and outcomes from complete relations. This establishes class-wise irredundancy within the declared failure model. Full constructions and projection assumptions are in the supplement.

3.4 Atomic successors and exact source preservation

For an admitted nonempty batch B with unique changed fields, the required effects are

$$\Delta(V_v, V_{v+1}) = \{i.f : i \in B\} = \{t.f : t \in T_{v+1}\}. \quad (7)$$

Each admitted item has one transition to its explicitly authorized value; all items share the target record, observed predecessor and one successor. Initial coverage establishes complete value/source domains. Later admissions preserve every unchanged field's value and exact source:

$$f \notin \Delta \Rightarrow V_{v+1}(f) =_{\text{json}} V_v(f), \quad S_{v+1}(f) = S_v(f). \quad (8)$$

Decision, authorization binding, transitions, complete version, sources and audit effects commit together, or the authoritative state stutters. Candidate and review preparation already persisted before the attempted admission remain historical evidence. The supplementary contract states no-op handling, mixed submissions and the declared field domain.

4 Formal and Transactional Realization

4.1 Making the relation executable

The relational model encodes candidates, evidence, principals, authorizations, ordered versions, complete states and source transitions. Machine events may extend candidates and evidence while framing authoritative state. Full admission adds binding, freshness, complete shape and item-transition correspondence to the base state change. Removing a conjunct while retaining another valid item exposes which required distinction that conjunct enforces. Alloy makes these finite relational checks executable [15, 16].

The archived checks cover legal admissions, conjunct-removal witnesses, source attacks and nonempty trace-preservation premises. All 72 command outcomes agree with their manifests across two scope profiles. The assertion checks have no counterexamples within their encoded finite scopes. An independent abstraction reads persisted pre/post relations, maps them into

fixed instances, and checks Full admission or rejected-state stutter: nine intended projections are satisfiable and twenty selected mapping mutants are unsatisfiable. The supplement records exact scopes, commands, assumptions and concrete/formal domain differences.

4.2 Transactional persistence and review binding

AUTO-DECTE separates candidate-writing capability from authoritative admission. An immutable certificate includes the persisted candidate instance and its canonical target, proposal, producer, expected predecessor and evidence identity. A review decision binds that certificate to the explicit authorized value. Field transitions retain the chain through both accepted and corrected proposals.

The relational adapter reloads the current field domain, complete predecessor, evidence and authorization within the transaction. It derives changed fields using canonical equality, validates exact bindings and unused identities, checks the principal policy, and compares-and-swaps the expected version. It then writes one transition per admitted field and copies exact unchanged sources into one complete successor. Constraints reinforce unique transitions and coherent references. Any unsuccessful attempt stutters relative to the immediately preceding authoritative database state.

An independently written event journal stores the same authority-relevant relations. Its retained-context configuration omits instance equality at the execution gate; its exact configuration adds the reviewed-candidate link. Comparing the exact journal with the relational implementation tests the same contract across storage representations. Comparing the two journal configurations isolates the binding choice within one representation.

Faithful transfer of review requires a binding before the final request arrives. A server-held review record can fix the candidate, expected version, reviewer and corrected value, then validate the request against that record. The resulting guarantee concerns request-to-record consistency. Browser presentation is a separately observable layer: a page can display a different candidate while the server and request still name the recorded one. Controlled request and display probes identify that boundary.

4.3 Independent audit of authoritative effects

The source verifier follows every field through its exact transition, authorization, candidate certificate, producer, evidence and predecessor. Changed fields need their unique new transition; unchanged fields retain the preceding source. A separate database oracle reads persisted relations rather than trusting the production planner or trace builder. Rejection checks compare full logical database digests, and corruption probes require changed state plus a diagnosis. These checks connect the written contract to the actual durable successor.

Table 1: Controlled mechanism evidence. Counts retain their original interpretation units; all cases are constructed or scripted.

Evidence	Observed result
Persistence binding	Exact journal/reference agree in 165 cases; context accepts 15 instance substitutions
Request binding	Server-held gate rejects 9 request substitutions accepted by original path
Display boundary	Both paths accept 3 display-only substitutions
Formal/concrete checks	72 bounded commands match; 9 intended projections SAT; 20 mapping mutants UNSAT

The prospective experiment uses another independent oracle for the paired policies. It checks raw responses against executed tool actions, hash-linked events, persisted state snapshots and reviewed-history provenance. Integrity is determined from authoritative transitions and audit coverage, rather than from a gate’s own claim that an operation was admissible. This permits a context-admitted transition to satisfy its policy while violating the separately declared instance-continuity oracle.

5 Controlled and Historical Evidence

5.1 Controlled mechanism and system validation

Context-equivalent substitutions separate reviewed-instance continuity from content approval. The archived E1 comparison executes 165 cases per mechanism. The context journal accepts all fifteen same-value/same-context substitutions; the exact journal and relational reference reject them. The two exact mechanisms agree on all corresponding decisions and provenance-query answer sets. Each accepts 45 legal cases and rejects 120 with authoritative-state stutter. This establishes the executable separator that the online experiment subsequently turns into agent-visible feedback.

Server-held review binding separates request substitution from faithful presentation. E2 has thirty browser cases per path: nine legitimate controls, nine request substitutions, nine stale/replay/interrupted controls and three display-only substitutions. The gate rejects every request substitution while preserving the legitimate controls. Both paths reject the stale, replay and interrupted cases without fact writes. Both admit the display-only controls because their actual request still matches the server record. Table 1 groups this evidence under one realization question.

Enforcement and tracing have measurable implementation costs. In the archived 128-field all-changed workload, median complete confirmation takes 712.00

ms in the relational reference and 22.09 ms in the exact journal. Their call boundaries include different connection and validation work. The separate prevalidated-materialization timer answers a different question and remains in the supplement. With the same trace verifier, the 128-field, 100-version, 1,000-record workload takes median 170.24 ms with bulk reads and 509.09 ms with point access, while full returned objects agree across all 36 paired cells. These costs characterize the evaluated persistence paths; agent-level friction is measured separately after policy intervention.

5.2 Observable correction in hosted workflows

The correction contract is recoverable from archived hosted-agent persistence. A new audit of 1,260 historical databases identifies 84 retained machine-origin correction chains, checking 252 evidence associations, 168 record versions and 168 reconstructed provenance traces. Each chain preserves a machine proposal, explicit host-authorized correction, exact candidate binding and persisted transition. These chains provide integration evidence for the role distinction defined in Section 3.

The audit unit is a persisted correction chain rather than whole-task success. The corresponding archived endpoints include 80 strict-safe task successes, one safe utility failure, two model/API failures and one timeout. All 84 persisted chains satisfy the audited relation. Their single-field cases provide zero unchanged-field checks, so copy-forward is supported by the controlled multi-field evidence. Full original task and annotation denominators remain in the supplement, separate from the prospective policy experiment.

6 Prospective Online Agent Experiment

Controlled mechanism separation establishes that the policies differ, but it does not tell us what happens once that difference becomes feedback to an adaptive agent. The prospective experiment compares Eqs. 4 and 5 in live continuations. It is designed around the sequence

$$\begin{aligned} &\text{feedback} \rightarrow \text{next agent action} \\ &\rightarrow \text{recovery path} \rightarrow (U, I). \end{aligned} \tag{9}$$

The local precollection freeze fixes the tasks, pair order, deployments, prompts, tools, resource rules, scorer and failure handling. The separately archived pilot supports engineering readiness and is excluded from formal effect analysis.

6.1 Shared review history and unique policy difference

A shared prospective prefix gives both policies the same reviewed history before execution binding becomes behaviorally relevant. An actual agent proposes and

obtains host review; the complete visible conversation and persisted state are then copied into isolated policy arms. The same task, checkpoint, tools, deterministic reviewer, persistence, prompt and resource-policy family apply to both arms. This isolates the effect of requiring I_{eq} on subsequent continuation. Both arms retain the reviewed identifier, so the comparison changes its use at execution rather than its availability for auditing.

A deterministic host reviewer fixes the correction decision across arms, separating authorization consequences from variation in review judgments. Two independently accessed deployments expose the same contract through distinct provider routes, with complete request-level provenance. Configuration A uses the frozen account-backed Codex route; B uses the recorded compatible API route. Requested aliases, observed returned identifiers, settings and collection timestamps are retained in the supplement and artifact. The common action interface presents the same tool definitions and returns actual results to subsequent decisions.

Each task/configuration forms one adjacent context/bound pair. A frozen assignment manifest balances which policy runs first within scenario and configuration. Two provider lanes execute one pair each, respecting this order. All planned arms remain in the analysis ledger, with runtime outcomes and audit coverage recorded separately. The supplement specifies the collection and interruption rules.

6.2 Task families and controlled exposure

The design contains 128 task instances, two configurations and two policies: 256 paired units and 512 planned policy trajectories, with one repetition. Controlled correction tasks expose specific authorization boundaries while keeping legal recovery paths interpretable. Their frozen distribution is

$$L32 + N32 + G16 + E16 + V16 + R16 = 128.$$

L: legal correction. Execute the reviewed instance with the corrected authorized value; check ordinary completion and nonvacuous unchanged-source preservation.

N: refreshed preview. After authorization, the agent obtains a required refreshed preview and decides how to preserve continuity. It can reuse the reviewed instance, authorize a new candidate or otherwise manage the workflow; the harness does not force an invalid commit.

G: queued handoff. The harness creates a same-value/same-context new instance and performs one standardized index action under the prior grant. Its actual result becomes feedback for autonomous continuation. The index action is marked `harness_handoff`; it is a standardized recovery probe.

E: changed evidence. Publish changed evidence bytes and locator while preserving the proposal value.

Both policies enforce the shared evidence/context checks; recovery can obtain a candidate reviewed against current evidence.

V: stale predecessor. Advance authoritative state externally. A fresh-state read, proposal and authorization provide a legal recovery path while preserving the externally updated source.

R: unavailable receipt. After a successful commit, the first receipt path is unavailable while authoritative state remains readable. Distinguish state confirmation, exact idempotent retry and semantic replay of a consumed grant.

The task generator varies business template, correction direction/magnitude, value relations, evidence layouts/locators, sources, version/history depths, retained metadata, candidate envelopes, snapshot organization, ordering and task wording. All 128 parameter vectors and rendered inputs are frozen. These variations diversify instances while preserving the six mechanisms.

6.3 Utility, integrity and recovery

Task utility $U = 1$ requires the authoritative goal, required workflow actions and an accurate final report. Integrity is independently scored as $I \in \{1, 0, \text{unknown}\}$. Covered authoritative transitions must satisfy authorization, freshness, evidence/context integrity, authorized-value equality, provenance, copy-forward and the declared instance-continuity requirement. A proved violation yields $I = 0$; incomplete necessary evidence yields unknown unless a violation is already proved. Runtime failure or rejected execution does not alone determine integrity.

Joint outcomes retain $(U, I) = (1, 1), (1, 0), (0, 1), (0, 0)$ and integrity unknown. A policy-admissible context substitution can therefore complete the task while failing the instance oracle. Completion and integrity are never collapsed into a total score.

A rejection is recoverable when its feedback is delivered and a legal path exists at the recorded rejection-time state. Delivery requires an actual received next model response with the recorded feedback input. Recoverability is determined from that state and available capabilities, not inferred backward from eventual success. The event-order classifier records the next meaningful tool action and observed reuse, verification, reproposal, reauthorization, state/evidence refresh, repeated invalid attempts, abandonment, confirmation and terminal failure. Raw component actions remain available alongside the primary path label. Every $U = 0$, $I = 0$ or unknown arm enters the complete failure-archaeology ledger.

6.4 Paired effects and operational components

Two primary endpoints are frozen: $\Delta U = U_{\text{bound}} - U_{\text{context}}$ and the bound-minus-context difference in run-level instance-continuity failure incidence. Transition-

level executed substitution is a separate diagnostic. Unknown continuity contributes explicit identification bounds rather than a zero replacement.

For aggregate effects, each task contributes its average configuration-specific paired difference. A scenario-stratified bootstrap resamples 128 task clusters, retaining both configurations and both policies; the frozen plan uses 10,000 iterations and seed 20261007. Separate configuration-specific exact McNemar tests are secondary. The structure supplies paired effects and controlled-task uncertainty, rather than 512 independent samples.

Recovery distributions conditional on delivered rejection are descriptive because exposure is post-treatment. G-bound provides the standardized probe cohort; N remains a separate account of agent-managed continuity. Operational components report additional tool calls, model turns, authorization attempts, reproposals, verifications, elapsed time and observed provider usage. Unsuccessful runs retain the specified budget-capped time-to-success alongside their actual termination time. Shared-prefix provider usage is charged once per pair. No arbitrary composite friction score is introduced.

7 Results

7.1 Continuity and task completion

Instance-bound execution separates accountable completion from successful state mutation in the standardized handoff tasks. The complete ledger contains 256 terminal pairs and all 512 planned arms. Context has 17 proved instance-continuity violations, all in G: sixteen in B and one in A. Bound has no proved violations; one bound arm and one context arm have unknown integrity. Fifteen context violations accompany successful task completion, and two accompany unsuccessful completion. The distinction between U and I therefore changes the interpretation of an apparently successful run.

The reviewed checkpoint is reached in 11 of A’s 128 pairs and all 128 of B’s pairs. In A, 117 prefixes terminate with `CODEX_EXIT_1`; the raw provider responses report an account usage limit. Both arms inherit that shared-prefix failure, yielding 234 unsuccessful planned arms with covered state and no unauthorized transition. Four further A suffixes terminate with runtime errors. B has eighteen malformed-output terminations across the two policies. The initial file-replacement interruption leaves one partial context arm and one unstarted bound arm with unknown integrity. These outcomes remain in Table 2, rather than being replaced with newly sampled tasks.

The paired completion estimate is small relative to its task-cluster uncertainty. Context completes 126/256 arms (49.22%) and bound 128/256 (50.00%). The frozen task-averaged bound-minus-context effect is +0.78 percentage points, with 95% interval $[-1.95, 3.52]$. In B, the corresponding counts are 117/128 and 119/128, an effect of +1.56 points with interval $[-3.91, 7.03]$. Five B tasks complete only under

context and seven only under bound; the secondary exact McNemar test gives $p = 0.774$. A has nine completions per policy. The intervals quantify variation over the paired controlled tasks and retain both deployments’ resource outcomes.

Unknown continuity outcomes produce an identification interval rather than a substituted point estimate. Across the complete planned ledger, the bound-minus-context difference in continuity-failure incidence lies in $[-7.03, -6.25]$ percentage points. The frozen bootstrap intervals for its lower and upper endpoints are $[-8.20, -6.25]$ and $[-7.42, -5.08]$, respectively. A single point estimate would conceal the two unknown arms. At the transition level, all seventeen observed substitutions are policy-admissible under context and instance-incompatible under the independent oracle. The remaining shared guard and provenance checks hold for these transitions. The improvement thus concerns the disclosed extra continuity obligation, rather than correcting a context policy’s own decision rule.

7.2 Feedback, next actions and recovery paths

Delivered instance-mismatch feedback redirects every exposed G-bound continuation toward an authorized execution path. The standardized handoff reaches seventeen of the 32 planned G pairs: sixteen in B and one in A. Each context index action accepts the new instance under the old reviewed grant. Each paired bound index action rejects it, and each rejection is proved delivered to an actual subsequent provider response. Legal recovery exists at all seventeen rejection states. All seventeen bound continuations end with $(U, I) = (1, 1)$, whereas the context continuations retain the irreversible instance violation even if later actions report the correct goal.

Reviewed-candidate reuse accounts for all sixteen B recoveries. Their first meaningful action after rejection is `commit` using the original reviewed candidate and grant. They neither regenerate evidence nor request a new authorization to solve this mismatch. A’s one delivered continuation instead requests authorization for the handoff candidate and then commits it. These are two distinct legal responses to the same continuity obligation. The counts describe actual observed responses: the host exposes the mismatch but does not prescribe either repair sequence.

Agent-managed refreshed previews preserve reviewed-instance continuity under both policies in every covered N transition. B has thirty such transitions per policy, all using an instance compatible with the review; A has three per policy. B completes 28 of its 32 N arms under each policy. Each B policy also has three covered unsuccessful arms and one unknown arm. No delivered instance-mismatch feedback occurs in N. Figure 3 keeps this absence of voluntary mismatch separate from G’s deliberately exposed handoff. The observed behavioral consequence of the unique policy conjunct is concentrated in G, rather than appearing throughout ordinary preview continuation.

Table 2: Complete planned-ledger outcomes. Every row contains 128 arms. A covered run without an unauthorized transition can have $I = 1$ despite unsuccessful execution. Unknowns are the two N arms affected by the initial infrastructure interruption.

Deployment	Policy	$U = 1, I = 1$	$U = 1, I = 0$	$U = 0, I = 1$	$U = 0, I = 0$	I unknown	$U = 1$	Checkpoint pairs
A	Context	8	1	119	0	0	9	11/128
A	Bound	9	0	119	0	0	9	11/128
B	Context	103	14	8	2	1	117	128/128
B	Bound	119	0	8	0	1	119	128/128




Figure 2: Task utility and integrity on the complete planned denominator. Deployment A’s large $U = 0, I = 1$ component records provider failures with covered state, predominantly before review. Deployment B supplies most online continuations. Successful context runs can contain an instance-incompatible transition even when the final authorized value is correct. Exact counts, including small cells, are in Table 2.

Evidence, version and receipt boundaries induce different observed recovery actions through the shared execution contract. In B, E supplies sixteen delivered evidence rejections per policy; fifteen continuations per policy obtain a fresh proposal and authorization and complete. V supplies thirteen delivered stale-version rejections per policy; twelve per policy recover through a current-predecessor proposal and review. The first V action varies between reading state and immediately proposing from available updated information. In R, nine context and ten bound receipt failures are proved delivered; nine continuations in each policy confirm success through state. One further bound run makes no subsequent tool action and still provides an accurate completed report, so the frozen episode classifier records no observed recovery endpoint for that episode. Run-level completion and the stricter episode recovery label remain separate.

The complete event-order analysis records observed strategies without filling absent categories with illustrative narratives. It identifies reuse, reproposal, reauthorization, state refresh and evidence refresh. No delivered episode is labeled reverify, repeated-invalid, fresh-id-invalid-loop, stale-reuse, hallucinated-success or timeout. Terminal malformed output and runtime interruption account for unsuccessful observed continuations; no frozen **abandon** label occurs. These episode distributions are descriptive conditional on actual feedback delivery. Table 3 therefore characterizes paths

through exposed boundaries rather than estimating an unconditional causal effect of rejection on recovery.

7.3 Operational friction and complete failure archaeology

The standardized continuity repair adds an explicit execution action even when task completion is preserved. Across B’s sixteen G pairs, bound uses exactly one additional agent tool call per pair, with no additional authorization requests. The median paired suffix-time difference is +14.54 seconds; its mean is +8.84 seconds. Model turns increase by a mean of 0.94. The one exposed A pair instead adds two tool calls, including one authorization request. The difference between reuse and reauthorization is therefore visible as an execution cost rather than merely as an extra rejection count.

Operational differences concentrate at the exposed boundary rather than uniformly increasing every component. Across B’s 127 complete suffix pairs, mean extra agent calls are 0.071, mean extra model turns 0.039 and mean extra authorizations zero; median extra elapsed time is 0.69 seconds. The 720-second budget-capped time-to-success endpoint retains all 128 B pairs and has mean paired difference −9.74 seconds. In G its mean is −76.53 seconds, despite positive median elapsed friction, because two unsuccessful context reports receive the unsuccessful-run cap while all bound runs complete. A favorable capped endpoint there-




Figure 3: Separating agent-managed continuity (N) from standardized mismatch recovery (G). N exposes a refreshed preview without forcing an invalid execution; G supplies a harness-origin index action whose actual acceptance or rejection becomes feedback. B’s 16 G-bound continuations reuse the reviewed candidate. A’s single exposed G-bound continuation obtains fresh authorization. Counts include the unobserved or unknown portions of each planned cohort.

Table 3: Delivered-feedback continuations in B. Each row’s denominator is delivered episodes, not all planned tasks. All listed episodes have a legal path at rejection time. E, V and R involve shared checks/capabilities; only G exposes the policies’ unique instance rule. Recovery requires a recorded recovery endpoint, $U = 1$ and $I = 1$.

Family	Policy	Delivered	Recovered	First meaningful action	Observed path
G	Bound	16	16	Commit (16)	Reuse reviewed candidate (16)
E	Context	16	15	Inspect evidence (15), propose (1)	Fresh proposal and authorization (15)
E	Bound	16	15	Inspect evidence (16)	Fresh proposal and authorization (15)
V	Context	13	12	Propose (7), state read (5), none (1)	Fresh proposal and authorization (12)
V	Bound	13	12	State read (10), propose (2), none (1)	Fresh proposal and authorization (12)
R	Context	9	9	State read (9)	State confirmation (9)
R	Bound	10	9	State read (9), none (1)	State confirmation (9)

Table 4: B’s paired operational components (bound minus context). “All” uses 127 pairs with both suffix records; G uses all 16 pairs. Elapsed time includes unsuccessful suffixes. Median and mean time answer different questions.

Component	All	G
Mean extra agent calls	0.071	1.000
Mean extra model turns	0.039	0.938
Mean extra authorizations	0.000	0.000
Median extra suffix seconds	0.69	14.54
Mean extra suffix seconds	1.29	8.84

fore does not imply faster execution. The supplementary component table retains actual time, capped time and provider usage separately, with the shared prefix charged once.

The all-arm archaeology contains 273 distinct arms with $U = 0$, $I = 0$ or unknown integrity. They par-

tition into 238 A runtime failures, fifteen successful context arms with an instance violation, two unsuccessful context arms with an instance violation, sixteen further unsuccessful B arms with covered intact integrity, and two unknown N arms. The unsuccessful B arms comprise eighteen malformed-output terminations overall, two of which also have the G continuity violation. Infrastructure interruption accounts for the remaining two unknown arms. Every violating transition is in G; no proved evidence, value, freshness, source-preservation or copy-forward violation accompanies it. The full ledger includes identities, stage, terminal type, component labels, missing coverage and the underlying episode records before any trajectory is selected for illustration.

The online evidence concerns controlled correction workflows and the two recorded deployment routes. Most live continuation evidence comes from B because A’s provider quota curtailed checkpoint coverage. G identifies recovery after a standardized mismatch, while

N supplies a separate observation of voluntary continuity management; neither cohort estimates the natural frequency of substitution. Together they establish where enforcement changed execution in this experiment and where it made no observed difference.

8 Discussion and Limitations

8.1 What continuity adds to corrected execution

Correction makes the authority for a value distinct from the origin of its proposal. Preserving only the final corrected value loses the machine output; preserving only the proposal loses the reviewer-authorized replacement. Retaining both still leaves the identity question: which candidate did the authorization cover, and which supplied the actual transition? The context-equivalent separator shows why those roles cannot be reconstructed from a successful endpoint alone. The fifteen $(U, I) = (1, 0)$ trajectories provide an online counterpart to that distinction: meeting the task goal and retaining the authorized value does not establish review-to-execution continuity.

This account gives established system mechanisms a common execution target. Transactions make the complete successor atomic. Provenance makes its effects reconstructible. Version validation controls the predecessor against which it is installed. Instance authorization determines which reviewed origin may supply that successor. Their joint contract supports the correction case without equating machine proposal and authorized value. The controlled agreement across relational and event-journal realizations shows that the relation can survive a change of representation; the hosted-chain audit shows that its roles can be recovered from persisted integration evidence.

8.2 Authorization as part of the execution environment

An authorization rule has behavioral consequences when its result becomes input to the next agent decision. In G, the difference between the policies is first mechanical: the same standardized request is accepted or rejected. The subsequent difference is an observed continuation. B responds by selecting the already reviewed candidate; A's available continuation obtains a new review. The gate does not generate those actions for the agent. It makes a continuity boundary visible, after which the agent selects an available legal path.

The contrast with N clarifies the mechanism's place in an agent system. A refreshed preview creates the possibility of identity substitution but does not force the agent to exploit it. Covered N transitions preserve the reviewed instance even when the looser policy would accept a context-equivalent alternative. The contract therefore performs two functions: it documents accountability when the agent already maintains continuity and supplies actionable feedback when a queued handoff breaks it. An execution-layer benefit can be

material even when voluntary mismatch is absent in the agent-managed cohort.

The broader recovery results also separate obligations from repair sequences. Changed evidence calls for a current evidence-based proposal and review; a stale predecessor calls for a current-state proposal; an unavailable receipt can be resolved by confirming state rather than applying another mutation. These paths use the same tool vocabulary but solve different contract boundaries. A useful execution environment should expose which relation failed and retain the state needed for a legal next step. Rejection frequency alone omits whether the agent received that information, acted on it and reached an accountable outcome.

8.3 Accountability and execution cost

The measured trade-off is additional execution work at a standardized continuity boundary, alongside improved instance accountability. The strongest direct result is B's G cohort: one extra agent call restores the reviewed link in every exposed pair, with a positive median suffix-time cost. Reauthorization is another legal strategy, but its one observed instance has a different call structure. These components matter operationally because a system can choose which recovery capabilities to expose and which authorizations remain reusable.

The overall completion interval includes both modest improvement and modest loss. It provides no basis for claiming utility equivalence or absence of completion cost. Equally, the capped time endpoint rewards accurately completed reports and should not be read as a latency speedup. Reporting separate calls, turns, review requests, elapsed time and U/I keeps these meanings visible. Provider failures additionally show why run completion, mutation integrity and availability need separate measurements: a covered run can leave state intact while accomplishing nothing.

8.4 Scope and further evaluation

The task family provides controlled exposure to six correction and execution boundaries, with one repetition of each task/configuration/policy combination. The deterministic reviewer fixes correction and authorization decisions, isolating execution binding from variation in human judgment; evaluations with interactive reviewers could additionally measure review burden and disagreement. The 128 tasks diversify inputs within those mechanisms, while broader business workflows, longer horizons and additional tools would test how the same contract interacts with a richer execution environment. The experiment measures controlled effects rather than production prevalence.

Deployment-level conclusions follow actual coverage. A reaches eleven checkpoints and contributes only one standardized mismatch recovery; B supplies the principal online evidence. Provider aliases and recorded settings support reconstruction of the accessed routes, while unavailable immutable model revisions constrain exact response reproduction. The frozen inputs, re-

ceived responses, checkpoints and independent scorer support retrospective verification of the reported transitions and paths. Additional prospectively frozen deployments would be needed to assess the stability of the observed strategy choices and completion costs across provider revisions and resource conditions.

9 Conclusion

Correction-aware review-to-execution continuity separates what an agent proposed, what a reviewer authorized, which candidate was reviewed and which supplied a persistent transition. Formal relations, transactional realization and independent auditing make that distinction executable across persistence representations. Controlled and historical evidence establish the mechanism and its observable integration roles.

The complete prospective ledger connects the contract to live continuation. Standardized handoffs expose seventeen instance violations under context authorization; bound execution rejects those handoffs and all seventeen exposed continuations complete with intact continuity through reviewed-candidate reuse or fresh authorization. Agent-managed preview tasks produce no observed substitutions. Recovery adds measurable execution work, and the paired completion uncertainty leaves the utility trade-off open. Review-to-execution continuity is therefore more than a property of stored records: enforced inside an agent loop, it becomes actionable feedback that can redirect execution and make the cost of accountability observable.

Statements and Declarations

Data and software availability

The study uses constructed tasks and public lifecycle inputs. The controlled experiment archive is fixed at commit 2645e5e18c900ea91c9c980e44195dc71e410432. The reference implementation, formal models and earlier historical evidence are retained at commit c6d512843c905cab6d8521dd8c914f7fb26d85ae. The prospective inputs and runner are fixed at master hash 7773563b40c8db30d6def121915f50fbad032537edde32afd4ba39ec3eea744b. The accompanying replication package retains the complete formal ledger, raw responses, state snapshots, scorer, tables, figure sources and file hashes; Online Resource 1 maps the layers and reproduction steps. Reanalysis requires no new model calls. New hosted runs require provider access.

Ethics approval and consent to participate

The prospective experiment uses constructed tasks and a deterministic host reviewer, with no new human participants. In the historical annotation study, one author (X.W.) and one non-author volunteer independently coded benign completion outcomes; L.H. adjudicated their disagreement and separately reviewed eleven

human-rule disagreements. X.W. also completed the textual-acknowledgment annotation. Only pseudonymous annotator identifiers were collected. Both coders were informed of the task and data use, and the non-author coder participated voluntarily. The authors assessed the annotation of generated outputs as not requiring ethics approval. No animal subjects were involved.

CRedit authorship contribution statement

Liang Hanghao: Conceptualization, Methodology, Software, Formal analysis, Investigation, Data curation, Validation, Visualization, Project administration, Writing—original draft, and Writing—review & editing. Xuan Wentao: Methodology, Formal analysis, Investigation, Data curation, Validation, Visualization, and Writing—review & editing. Chen Qile: Investigation and Validation. Peng Peng: Supervision and Writing—review & editing.

Funding and competing interests

This research received no external funding. The authors declare no competing interests.

AI-assisted preparation

DeepSeek, OpenAI Codex and Claude assisted with code, analysis checks, literature organization and manuscript preparation. OpenAI Codex assisted with the present whole-paper reconstruction and figure scripts. An earlier diagram prototype used OpenAI image generation with an unrecorded model version; the present figures are rendered from explicit plotting scripts. AI tools are not authors. The authors retain responsibility for the article and its evidence.

Supplementary information

Online Resource 1 contains formal constructions and checks, full controlled workloads, historical records, prospective task and deployment definitions, scoring and interruption rules, complete outcome and recovery tables, and the artifact map.

References

- [1] LangChain. Human-in-the-loop. Official documentation. <https://docs.langchain.com/oss/python/langchain/human-in-the-loop>. Accessed 7 October 2026.
- [2] Yakout M, Elmagarmid AK, Neville J, Ouzzani M, Ilyas IF (2011) Guided data repair. *Proc VLDB Endow* 4(5):279–289. <https://doi.org/10.14778/1952376.1952378>.
- [3] He J, Veltri E, Santoro D, Li G, Mecca G, Papotti P, Tang N (2016) Interactive and Deterministic

- Data Cleaning. Proc SIGMOD, pp 893–907. <http://doi.org/10.1145/2882903.2915242>.
- [4] Gray J (1981) The Transaction Concept: Virtues and Limitations. Proc Seventh International Conference on Very Large Data Bases, pp 144–154.
 - [5] Kung HT, Robinson JT (1981) On optimistic methods for concurrency control. ACM Trans Database Syst 6(2):213–226. <https://doi.org/10.1145/319566.319567>.
 - [6] Buneman P, Chapman A, Cheney J (2006) Provenance management in curated databases. Proc SIGMOD, pp 539–550. <https://doi.org/10.1145/1142473.1142534>.
 - [7] Moreau L, Missier P (eds) (2013) PROV-DM: The PROV Data Model. W3C Recommendation. <http://www.w3.org/TR/2013/REC-prov-dm-20130430/>.
 - [8] Clark DD, Wilson DR (1987) A Comparison of Commercial and Military Computer Security Policies. IEEE Symposium on Security and Privacy, pp 184–194. <https://doi.org/10.1109/SP.1987.10001>.
 - [9] OWASP Foundation. Transaction Authorization Cheat Sheet. https://cheatsheetseries.owasp.org/cheatsheets/Transaction_Authorization_Cheat_Sheet.html. Accessed 7 October 2026.
 - [10] Mershad K, Malluhi QM, Ouzzani M, Tang M, Aref WG (2015) Approving Updates in Collaborative Databases. IEEE International Conference on Cloud Engineering, pp 42–47. <https://doi.org/10.1109/IC2E.2015.31>.
 - [11] Buneman P, Khanna S, Tan WC (2001) Why and Where: A Characterization of Data Provenance. ICDT, LNCS 1973, pp 316–330. https://doi.org/10.1007/3-540-44503-X_20.
 - [12] Cheney J, Chiticariu L, Tan WC (2009) Provenance in Databases: Why, How, and Where. Found Trends Databases 1(4):379–474. <https://doi.org/10.1561/19000000006>.
 - [13] Arab BS, Gawlick D, Krishnaswamy V, Radhakrishnan V, Glavic B (2018) Using Reenactment to Retroactively Capture Provenance for Transactions. IEEE Trans Knowl Data Eng 30(3):599–612. <https://doi.org/10.1109/TKDE.2017.2769056>.
 - [14] Debenedetti E, Zhang J, Balunovic M, Beurer-Kellner L, Fischer M, Tramèr F (2024) AgentDojo: A Dynamic Environment to Evaluate Prompt Injection Attacks and Defenses for LLM Agents. Advances in Neural Information Processing Systems 37:82895–82920. <https://doi.org/10.52202/079017-2636>.
 - [15] Jackson D (2002) Alloy: a lightweight object modelling notation. ACM Trans Softw Eng Methodol 11(2):256–290. <https://doi.org/10.1145/505145.505149>.
 - [16] Torlak E, Jackson D (2007) Kodkod: A Relational Model Finder. TACAS, LNCS 4424, pp 632–647. https://doi.org/10.1007/978-3-540-71209-1_49.
 - [17] Marinov D, Khurshid S (2001) TestEra: A Novel Framework for Automated Testing of Java Programs. IEEE International Conference on Automated Software Engineering, pp 22–31. <https://doi.org/10.1109/ASE.2001.989787>.