

Auto-Decte: An LLM- and RAG-Assisted Human-in-the-Loop Document Intelligence System for Industrial Paper Forms

Anonymous Authors

Hunan University, Changsha, China
(anonymized for review)

Abstract. Deploying large language models (LLMs) in high-stakes enterprise workflows is constrained less by model capability than by trust: suggestions must never silently become facts, every decision must be traceable to evidence, and the system must keep working when the model is unavailable. This paper presents Auto-Decte, a document intelligence system for industrial paper forms that operationalizes these constraints through a human-in-the-loop architecture. A fiducial-guided perception front end (QR/ArUco template recognition, perspective correction, per-cell digit and OMR recognition with an explicit ambiguity band) converts photographed paper forms into structured candidate values. An LLM suggestion adapter (DeepSeek, structured output) and a local vector similarity-retrieval module (RAG-style, read-only) assist reviewers, while a React review workbench makes the human the only source of truth: machine outputs are stored as append-only candidates, and only human confirmations or corrections become versioned, auditable records. A PWA front end with an offline outbox extends submission to in-plant mobile devices. The system is evaluated on end-to-end acceptance (389 public-snapshot tests, including the AI-off path and suggestion-isolation invariants), recognition benchmarks, a baseline comparison against a learned recognizer, end-to-end whole-form cell accuracy (97.5% clean), and per-image latency.

Keywords: LLM application · retrieval-augmented generation · human-in-the-loop · document intelligence · traceability

1 Introduction

Foundation-model-assisted document processing is attractive for the long tail of paper-based workflows in small and medium plants, but its adoption in payroll-adjacent settings is blocked by a trust problem: an LLM that transcribes a digit wrong changes a worker’s wage, and a retrieval system that injects a similar-looking historical record into a decision can silently corrupt it. Fully automatic pipelines are therefore unacceptable where errors carry financial consequences, yet purely manual digitization does not scale.

Auto-Decte addresses this with a human-in-the-loop architecture built on three principles: (i) *machine outputs are candidates, never facts* — the LLM adapter and every recognition stage emit candidates that a human must confirm or correct; (ii) *the document is engineered* — templates are generated with QR codes and ArUco fiducials so that perception is reliable rather than learned; (iii) *everything is auditable* — evidence is content-addressed, records are append-only versions, and every exported cell is back-traceable. The system is deployed as a working baseline at a bamboo-processing plant.

Our contributions are:

1. A human-in-the-loop document intelligence architecture integrating a fiducial perception front end, an LLM suggestion adapter, and local similarity retrieval under a single “candidates-only, human-confirmed” trust boundary.
2. Trustworthy LLM/RAG integration mechanisms: suggestions are stored separately and cannot mutate confirmed values; retrieval results are strictly read-only; the core pipeline remains correct with the AI layer fully disabled.
3. An engineered, interpretable perception front end whose explicit ambiguity band routes borderline cells to humans instead of guessing.
4. An offline-capable mobile submission path (PWA + IndexedDB outbox + idempotency keys).
5. A deployment-motivated case study with end-to-end acceptance and honest reporting of what remains unmeasured.

2 Related Work

Document OCR and layout analysis. General document AI stacks (PP-OCRv3 [1], LayoutLMv3 [2], Donut [3]) target free-form documents with learned detection and recognition. Auto-Decte deliberately keeps learned layout models out of the critical path: system-generated templates make layout known exactly, so perception reduces to fiducial-guided cropping plus per-cell template matching; learned OCR remains future work for handwriting and Chinese text.

LLM assistance and retrieval. Retrieval-augmented generation and generative information-extraction surveys [8, 9] and local-first software principles [10] motivate a read-only DeepSeek structured-output suggestion adapter and a local vector similarity index. The distinction from generic RAG is architectural: in Auto-Decte the model and the retriever are denied write access to facts by construction.

Human-in-the-loop digitization. HITL principles for high-stakes data [7] motivate the review workbench: machine outputs are append-only candidates, and every confirmed or corrected value is a versioned human decision with a before/after audit pair.

Fiducial-based localization and OMR. ArUco detection [5] and MILP marker dictionaries [6] provide geometric anchoring; existing OMR systems (e.g., SurveyNet [4]) evaluate bulk mark recognition. Auto-Decte’s distinguishing feature is the first-class ambiguity band: borderline cells are emitted with zero confidence and force review rather than a thresholded guess.




Fig. 1. End-to-end pipeline. Recognition and the LLM produce candidates and references; only human confirmations or corrections become facts.

Positioning. Four decisions distinguish Auto-Decte from prior systems: (i) machine outputs — from recognition and from the LLM alike — are candidates, never facts; (ii) an explicit ambiguity band routes borderline cells to humans; (iii) full traceability; (iv) offline mobile submission integrated with the same review queue.

3 System Overview

Auto-Decte is a single-machine deployment: FastAPI (Python 3.11) backend, SQLite via SQLAlchemy 2.0 with Alembic migrations, React 18 + Vite front end, OpenCV for image operations, a DeepSeek client for structured-output suggestions, and a local vector index for similarity retrieval. Fig. 1 shows the pipeline.

The domain model separates immutable evidence, versioned form records, append-only recognition attempts, isolated AI suggestions, read-only retrieval references, audit events, and export batches. Roles (operator, reviewer, finance, admin, auditor) gate every mutation.

4 Template Design and Perception Front End

Templates are designed and published inside the system (template center with versioning and draft canvas). Each template renders a PDF/PNG with a QR code (identity + version) and ArUco markers at known positions. Recognition is fiducial-first: a multi-region QR search classifies the form, four ArUco corner markers (IDs 10–13) map it to the canonical canvas, perspective correction aligns it, and per-field cells are cropped from the template definition.

5 Field-Level Recognition with Confidence and Ambiguity

Digit cells. Each cell is normalized to 48×64 , binarized with Otsu thresholding, and matched against ten rendered synthetic-font templates (0–9) by normalized mean absolute difference. The recognizer reports the best digit, a distance-based confidence, and an **AMBIGUOUS** flag when the margin to the second-best template falls below 0.02. Blank cells (ink ratio $< 1\%$) are reported as blank rather than as a digit.

Checkbox (OMR) cells. Fill ratio in the cell interior is computed after Otsu binarization; ratios ≤ 0.15 map to unchecked, ≥ 0.35 to checked, and the $(0.15, 0.35)$ band maps to **AMBIGUOUS** with zero confidence. The ambiguity band routes exactly the cells that recognition cannot decide to the human, instead of silently guessing.

6 LLM Assistance and Similarity Retrieval

Two optional AI components support the reviewer, both unable to write facts: (i) *LLM suggestions* — a DeepSeek adapter with structured output requests field suggestions; suggestions are stored separately, carry their evidence reference, and cannot mutate confirmed values; the adapter is disabled by default so the core pipeline runs without any external model (an acceptance-tested AI-off path). (ii) *Similarity retrieval* — a local vector index shows similar historical forms as reference; retrieval results are strictly read-only. The trust boundary is architectural: the model and the retriever produce candidates and references; only the reviewer produces facts.

7 Data Integrity, Versioning, and Export

Evidence images are content-addressed with SHA-256; duplicates are detected and never silently overwritten. Every human correction appends a versioned record with before/after audit rows. XLSX export produces four worksheets (formal data, exceptions & review, summary, export notes) and every exported cell is back-traceable to form, version, evidence image, and event log.

8 Offline Mobile Submission

A PWA front end supports in-plant phone use: credentials are salted with crypt; sessions use HttpOnly SameSite cookies with CSRF double-submit; drafts live in an IndexedDB outbox keyed by owner/device with idempotency keys; retries back off on network errors and stop on 4xx conflicts; mobile submissions flow into the same review path as scanned forms.

Table 1. Digit recognition: template matcher vs learned HOG+SVM baseline (per font, averaged over 7 perturbation conditions).

Font	SVM acc.	TM guess	TM confirmed	TM ambiguous
Simplex (printed)	1.000	1.000	1.000	0.000
Duplex (printed)	1.000	1.000	1.000	0.000
Triplex (unseen)	0.803	0.863	0.754	0.146
Complex (unseen)	0.797	0.894	0.797	0.103

9 Evaluation

Methodology: end-to-end functional acceptance via the automated test suite; controlled benchmarks on the perception front end using the system’s canonical template conventions and synthetic perturbations (blur, rotation, brightness, Gaussian noise, perspective tilt); fill-ratio sweeps for OMR; per-stage latency on A4@150 dpi images.

Functional acceptance and trustworthy-AI invariants. On the public snapshot of the integration branch, 389 tests pass (import, SHA-256 dedup, QR classification, digit/OMR candidates, rule validation, human confirmation, export, reverse traceability), and the suite enforces the AI trust boundaries: the AI-off path completes the workflow end-to-end; LLM suggestions are stored separately from confirmed values; similarity-retrieval results are read-only. 159 additional tests require enterprise database fixtures excluded from the public repository by data policy.

Digit recognition and baseline comparison. Table 1 reports the template matcher against a learned HOG + linear-SVM recognizer trained on the printed font. On the printed font family both reach 100%; on unseen fonts they are comparable, but only the template matcher routes 10–15% of cells to human review instead of silently guessing. Against a general OCR engine (Tesseract, single-character mode) on the same test set, the template matcher wins on the deployment-relevant printed-template regime (100% vs 84–87%), whereas Tesseract generalizes better to unseen fonts (94–95% vs 86–90%), which is outside the printed-form scope; in all cases only the template matcher emits an ambiguity channel.

OMR. Fill-ratio sweep: ratios ≤ 0.15 classified unchecked (100% correct), ≥ 0.50 classified checked (100% correct), and (0.15,0.50) emitted ambiguous (100% routed to review).

Template detection and geometry. QR payload recovery succeeds on 100% of 90 perturbed renders (blur, noise, brightness, $\pm 5^\circ$ rotation, 8% tilt). ArUco-based perspective correction succeeds on 100% of clean/blur/noise/brightness/tilt images with mean alignment error ≤ 0.34 px, dropping to 50%/20% at $3^\circ/5^\circ$ rotation — the known weak axis of fiducial recovery.

End-to-end whole-form evaluation. Synthetic filled forms (10 digit cells with random printed fonts + 10 OMR boxes) pass the full pipeline. Table 2 reports cell-level accuracy: 97.5% clean, 96.5–98.0% under blur/noise/brightness/tilt,

Table 2. End-to-end whole-form cell-level accuracy (synthetic filled forms, 20 forms per condition).

Perturbation	Digit acc.	Digit ambiguous	OMR acc.	Form corrected
Clean	0.975	0.080	0.940	1.000
Blur 3×3	0.965	0.055	1.000	1.000
Noise $\sigma=20$	0.970	0.085	0.970	1.000
Brightness $0.7\times$	0.980	0.045	0.965	1.000
Tilt 8%	0.980	0.080	1.000	1.000
Rotation 3°	0.575	0.025	0.600	0.600

with 4.5–8.5% of digit cells routed to review; degradation appears only under rotation.

Latency. Mean per-image stage times: QR search 2,180 ms (multi-region robust search dominates), quality assessment 20 ms, perspective correction 35 ms.

Similarity retrieval. The read-only retrieval module (character-bigram + word cosine over form records) is evaluated on a synthetic corpus of 200 records across five templates, with relevance defined as sharing the template key (the reviewer’s notion of a “similar past form”). Mean precision is 1.00@1, 0.99@3, 0.975@5, and 0.95@10, confirming that the module surfaces the intended same-template references without injecting them into records.

LLM/RAG effectiveness. The AI layer’s correctness and isolation are established functionally (AI-off path, suggestion isolation, read-only retrieval); the retrieval module is quantitatively validated (Precision@k above), while the end-to-end efficiency gain of LLM suggestions remains a field-pilot future work.

10 Limitations and Future Work

General handwritten digit-string / Chinese text / signature OCR is not implemented (only single-cell template digits and OMR); deployment is single-machine without a completed multi-user capacity test or field pilot; the evaluation lacks a 30–50-form golden set with double entry and real plant images, so accuracy claims are limited to controlled benchmarks; and the LLM suggestion module has functional correctness but no measured end-to-end efficiency gain. Future work: general OCR integration, independent task workers, external storage, a quantitative LLM-assistance study, and a field pilot measuring financial-processing time reduction.

11 Conclusion

Auto-Decte demonstrates a trust-first integration of an LLM adapter and similarity retrieval into an industrial document workflow: recognition and the model produce candidates and references, while humans remain the only source of facts, with full traceability from photograph to exported cell. The system is deployed

as a working baseline at a bamboo-processing plant and provides the substrate for a future quantitative field study.

References

1. C. Li et al.: PP-OCRv3: More attempts for the improvement of ultra lightweight OCR system. arXiv:2206.03001 (2022)
2. Y. Huang et al.: LayoutLMv3: Pre-training for document AI with unified text and image masking. arXiv:2204.08387 (2022)
3. G. Kim et al.: OCR-free document understanding transformer. arXiv:2111.15664 (2022)
4. SurveyNet: A unified deep learning framework for OCR and OMR-based survey digitization. *Journal of Imaging* 12(4), 175 (2026). DOI 10.3390/jimaging12040175
5. F.J. Romero-Ramirez, R. Muñoz-Salinas, R. Medina-Carnicer: Speeded up detection of squared fiducial markers. *Image and Vision Computing* 76, 38–47 (2018)
6. S. Garrido-Jurado, R. Muñoz-Salinas, F.J. Madrid-Cuevas, M.J. Marín-Jiménez: Generation of fiducial marker dictionaries using mixed integer linear programming. *Pattern Recognition* 51, 481–491 (2016)
7. S. Budd, E.C. Robinson, B. Kainz: A survey on active learning and human-in-the-loop deep learning for medical image analysis. *Medical Image Analysis* 71, 102062 (2021)
8. D. Xu et al.: Large language models for generative information extraction: A survey. arXiv:2312.17617 (2023)
9. Y. Gao et al.: Retrieval-augmented generation for large language models: A survey. arXiv:2409.14924 (2024)
10. M. Kleppmann, A. Wiggins, P. van Hardenberg, M. McGranaghan: Local-first software: You own your data, in spite of the cloud. In: *Onward!* (2019). DOI 10.1145/3359591.3359737